

FREE-ENERGY PRINCIPLE (FEP) TUTORIAL

Perception as an inference process - predictive processing

Author: [Pablo Lanillos](#) Assistant Professor Cognitive AI. Donders Institute for Brain, Cognition and Behaviour. The Netherlands.

Adapted to NeurIPS lecture by: Marius Marten Kästingschäfer, Marcel van Gerven

Background reading

Papers:


- Bogacz, R. (2017). A tutorial on the free-energy framework for modelling perception and learning. *Journal of mathematical psychology*, 76, 198-211.
- Lanillos, P., & Cheng, G. (2018, October). Adaptive robot body learning and estimation through predictive coding. In 2018 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS) (pp. 4083-4090). IEEE.
- Sancaktar, C., van Gerven, M., & Lanillos, P. (2019). End-to-end pixel-based deep active inference for body perception and action. *arXiv preprint arXiv:2001.05847*.
- Friston, K. J., Daunizeau, J., Kilner, J., & Kiebel, S. J. (2010). Action and behavior: a free-energy formulation. *Biological cybernetics*, 102(3), 227-260.
- Buckley, C. L., Kim, C. S., McGregor, S., & Seth, A. K. (2017). The free energy principle for action and perception: A mathematical review. *Journal of Mathematical Psychology*, 81, 55-79.

Blog:

- <https://msrmblog.github.io/is-this-my-body-1/>
- <https://msrmblog.github.io/is-this-my-body-2/>

Problem definition

We have an agent that lives in a 1-dimensional space environment. The agent goal is to infer its position in the space x given the sensory input s . It has a sensor that it is able to sense its rotation over the environment. After exploring the 1-dimensional space, it has learnt a generative model of the world depending on its internal state x , which is x^2 - See image below. *This generative model defines the prediction of the sensory input depending on its internal state.*



The agent starts with a wrong belief of his internal state, which yields into a bad estimation of its position. As it can be seen in the image below his actual position is misaligned with his estimated position. We will use a simplified model of predictive processing minimizing the Variational Free Energy (VFE) to solve this task.



Overview

In this toy example you will see a basic FEP model at work. The following will be done:

- Look at the formulas for the variational free energy
- Implement the generative model
- Run the generative model
- Inspect the results

Tasks:

- Implement the [updating step](#) for the generative model
- Test its robustness in regard to [different parameter](#) initializations
 - What effects does it have if the initial values are further away from the agents actual position?
 - What values will make the system to be stable?
 - What if we initialize the belief with a negative position?
 - What happens if we add noise to the sensory measurement?
- Describe according to the FEP how the agent would [act](#) in this environment

Part a: Variational Free Energy (VFE) - with the Laplace approximation

$$F = KL[q(x)||p(x|s)] - \ln p(s) \rightarrow_{\text{laplace}} F = -\ln p(s, \mu) - \frac{1}{2}(\ln |\Sigma| + \ln 2\pi)$$

Perception is defined as the minimization of F :

$$\mu = \arg \min_s F(\mu, s)$$

We compute it using the following differential equation (gradient optimization):

$$\dot{\mu} = -\frac{\partial F}{\partial \mu} = -\frac{\partial \ln p(s|\mu)p(\mu)}{\partial \mu}$$

(Eq. 1)

Assuming that the sensors measurements s are defined by a generative model with Gaussian noise z : $s = g(\mu) + z$

We have:

$$\dot{\mu} = \frac{\partial g(\mu)}{\partial \mu}^T \Sigma_s^{-1} (s - g(\mu))$$

(Eq. 2)

Question a.1: Write down the derivation from (Eq. 1) to (Eq. 2)

1. The likelihood $p(s|\mu)$ that defines the generative model follows a Normal distribution with mean μ and σ standard deviation: $p(s|\mu) \sim \mathcal{N}(s; \mu, \sigma^2)$
2. The prior $p(\mu)$ is assumed to be uniform

Tip: use the function of the Normal density function (pdf). [wiki](#)

Your derivation here:

$$\begin{aligned}\dot{\mu} &= -\frac{\partial F}{\partial \mu} = -\frac{\partial \ln p(s|\mu)p(\mu)}{\partial \mu} \\ &= -\frac{\partial}{\partial \mu} \ln\left(\frac{1}{\sqrt{2\pi\Sigma}} \exp\left(\frac{1}{2} \frac{(s - g(\mu))^2}{\Sigma}\right)\right) \\ &= -\frac{\partial}{\partial \mu} \ln\left(\frac{1}{\sqrt{2\pi\Sigma}}\right) + \ln\left(\exp\left(\frac{1}{2} \frac{(s - g(\mu))^2}{\Sigma}\right)\right)\end{aligned}$$

The $\ln\left(\frac{1}{\sqrt{2\pi\Sigma}}\right)$ factor equals to 0 as it doesn't contain a μ term. And the \ln and \exp also cancel each other out.

$$\begin{aligned}&= -\frac{\partial}{\partial \mu} \frac{1}{2} \frac{(s - g(\mu))^2}{\Sigma} \\ &= -\frac{\partial}{\partial \mu} \frac{1}{2} \Sigma^{-1} (s - g(\mu))^2 \\ &= -\frac{1}{2} \Sigma^{-1} \frac{\partial}{\partial \mu} (s - g(\mu))^2\end{aligned}$$

Apply the derivative rule

$$= -\frac{1}{2} \Sigma^{-1} 2(s - g(\mu)) \frac{\partial}{\partial \mu} g(\mu)$$

2's cancell out and rewriting gives...

$$\dot{\mu} = \frac{\partial g(\mu)^T}{\partial \mu} \Sigma_s^{-1} (s - g(\mu))$$

Part b: Generative model definition

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

# Generative model definition
def g(x):    # forward model
    return x**2

def dg(x):   # inverse model
    return 2*x
```

```
In [2]: # Plotting actual and inferred position of the agent
```

```
def plot_results(real_x, mu_0):

    x = np.linspace(-2, 2, 1000)
    y = np.linspace(2, 2, 1000)
    plt.plot(x, y, 'k')
    plt.plot(real_x, 2, 'r8', markersize=15)
    plt.plot(mu_0, 2, 'b8', markersize=15)
    plt.plot(0, 2, 'k|', markersize=8)

    plt.xlim([-1,1]) # adjust plot when choosing values outside of range -1 and 1
    plt.legend(['Space', 'actual position', 'estimated position', 'x=0'])
    plt.xticks([])
    plt.yticks([])
    plt.title('Estimated vs. actual position')

    plt.show()
```

Initialization

```
In [3]: dt = 0.001 # integration step
        ftime = 1 # maximum time considered

        time = np.arange(0, ftime, dt) # Time array
        N = time.shape[0] # Steps

        real_x = 0.1 # real position of the ball
        real_x_dot = 0 # change on the position

        mu_0 = 0.5 # internal belief
        mu_dot = 0.0 # change on the internal belief

        sigma_s = 0.01 # variance of the sensory process
        z = 0 # for now
```

Question b.1: Compute the sensor measurement

```
In [4]: # your code here

        s_measure = g(real_x) + z
        s_measure
```

Out[4]: 0.010000000000000002

Question.b.2: Compute the predicted sensor value depending on the brain variable μ

```
In [5]: # your code here

        mu_dot = dg(mu_0) * sigma_s ** -1 * ( s_measure - g(mu_0))
        mu_0 += dt * mu_dot

        mu_0
```

Out[5]: 0.476

Question b.3: Compute the prediction error

In [6]:

```
# your code here

e_s = s_measure - mu_0
e_s
```

Out[6]: -0.46599999999999997

Part c: Implementing the FEP model

Question c.1: Write the function fep(parameters) that executes the FEP model

In [7]:

```
def fep(real_x, mu_0):
    """
    Takes real position real_x and belief in position mu_0 as input and runs FEP for
    Returns data which is a 4 x N array that collects s, real_x, mu_0 and e_s for N
    """
    data = np.zeros((4, N)) # (4, 1000)

    for i in range(0, N):
        s = g(real_x) # sensor reading
        mu_0 += dt * (dg(mu_0) * sigma_s ** -1 * (s - g(mu_0))) # change in mu_0
        e_s = s - mu_0 # error between sensor and belief

        data[0,i] = s
        data[1,i] = real_x
        data[2,i] = mu_0
        data[3,i] = e_s

    return data

# Initialization
real_x = 0.1 # real position of the ball
mu_0 = 0.5 # internal belief

data = fep(real_x, mu_0)
```

Part d: Inspect the results

In [8]:

```
def plot_variables(data):

    # Plot variables

    plt.figure(figsize=(10, 10))
    plt.subplots_adjust(hspace=0.3)

    ax1 = plt.subplot(311)
```

```

ax1.plot( time, data[0, :] )
ax1.legend(['s'])
ax1.title.set_text('Sensory measurment')


ax2 = plt.subplot(312)
ax2.plot(time, data[1, :], 'b' )
ax2.plot(time, data[2, :], 'r' )
ax2.legend(['real_x', 'mu_0'])
ax2.title.set_text('Actual vs. assumed position')

ax3 = plt.subplot(313)
ax3.plot(time, data[3, :] , 'k' )
ax3.legend(['e_s'])
ax3.set_xlabel('time')
ax3.title.set_text('Prediction error')

plt.show();

plot_variables(data)

```



Question d.1: Now test the same with the following agent initialization and explain what happened


In [9]:

```

real_x = 0.1    # real position of the ball
mu_0 = -0.5     # internal belief

data = fep(real_x, mu_0)
plot_variables(data)

```



Describe your findings here

We see that our sensor reading is constant, and that if our initial internal belief is positive that we converge to 0.1. However, if our initial internal belief is negative (-0.5) then we only converge to -0.1. This most likely has something to do with the models **g** and **dg**.

Question d.2: Now add different levels of Gaussian noise to the sensor measurements and run the FEP model again

Write a function `fep2` with an extra parameter that is the measurement standard deviation σ , and plot the variables

$s = g(x) + N(0, \sigma^2)$ - Increase σ to increase the input noise

```

In [12]: def fep2(real_x, mu_0, measurement_std, sigma_s):
          """

```

Takes real position `real_x`, belief in position `mu_0` and standard deviation `sigma`

Returns data which is a 4 x N array that collects `s`, `real_x`, `mu_0` and `e_s` for N

"""

```
data = np.zeros((4, N)) # (4, 1000)
```

```
for i in range(0, N):
```

```
    s = g(real_x) + np.random.normal(0, measurement_std**2)
```

```
    mu_0 += dt * (dg(mu_0) * sigma_s ** -1 * (s - g(mu_0)))
```

```
    e_s = s - mu_0
```

```
    data[0,i] = s
```

```
    data[1,i] = real_x
```

```
    data[2,i] = mu_0
```

```
    data[3,i] = e_s
```

```
return data
```

```
real_x = 0.1 # real position of the ball
```

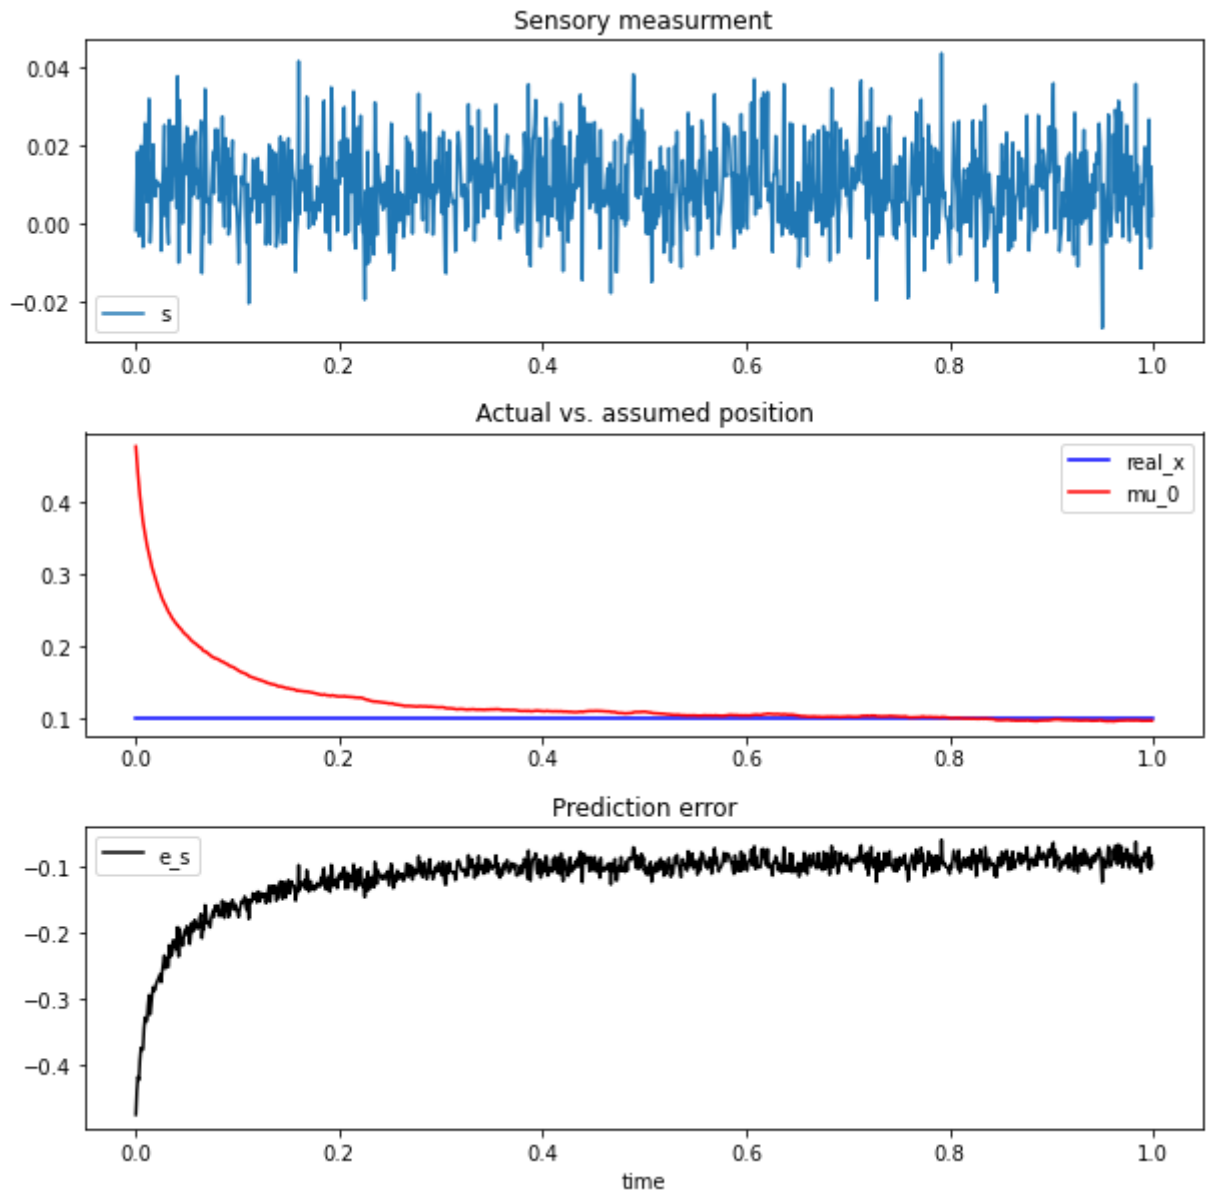
```
mu_0 = 0.5 # internal belief
```

```
measurement_std = 0.1
```

```
sigma_s = 0.01 # variance of the sensory process this is for the next exercise
```

```
data = fep2(real_x, mu_0, measurement_std, sigma_s)
```

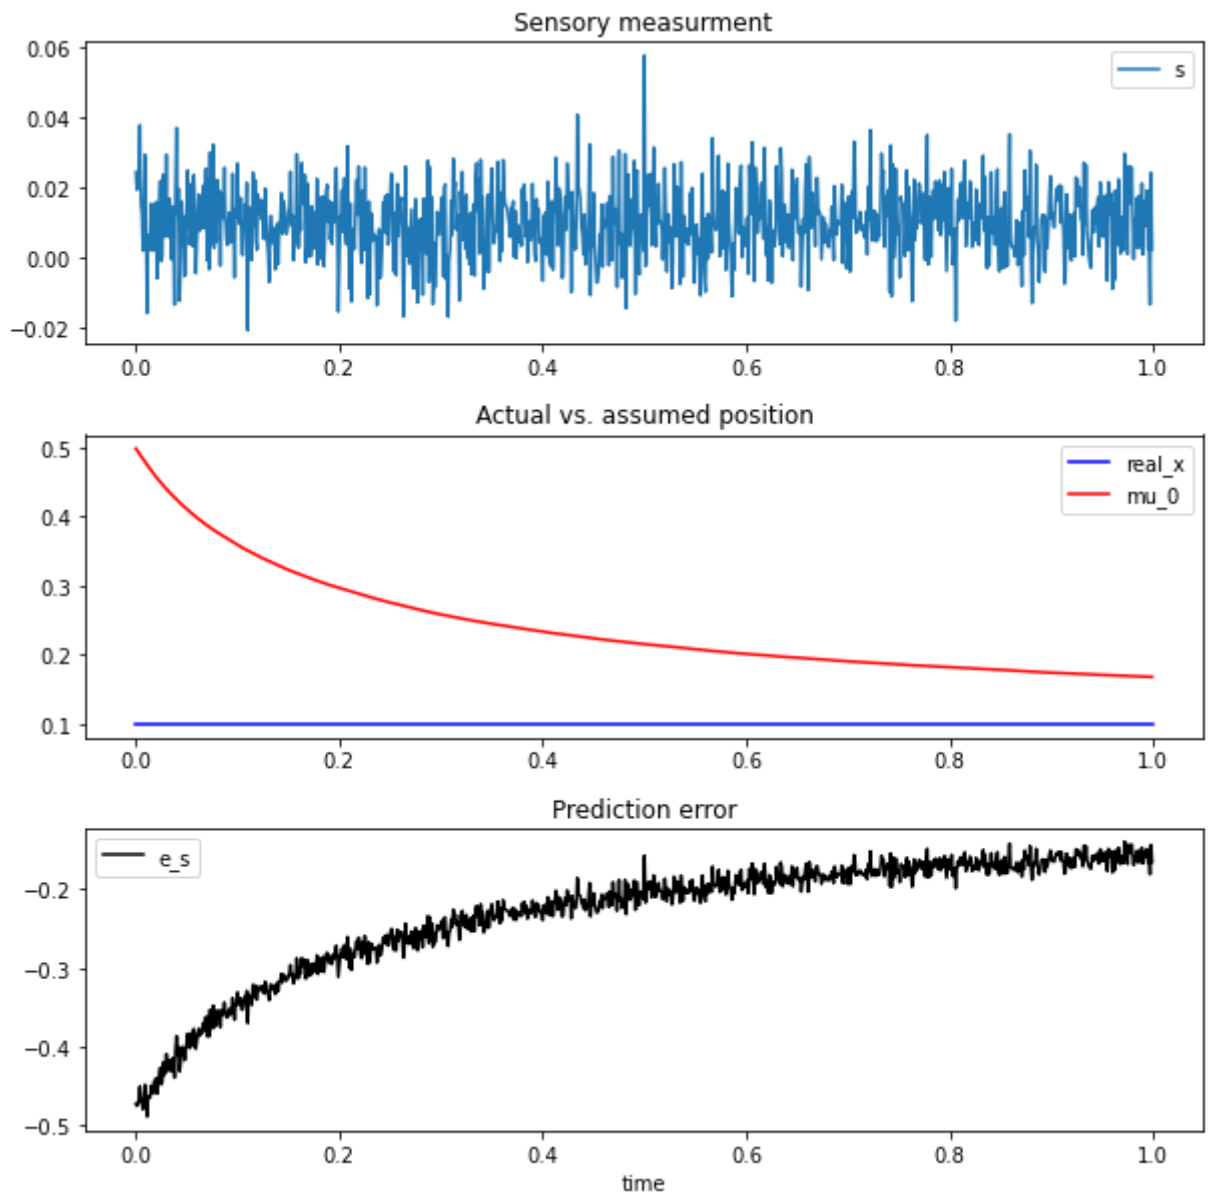
```
plot_variables(data)
```

Question d.3: Modify the variance parameter σ that weights the sensor error accordingly and run it again

```
In [11]: real_x = 0.1    # real position of the ball
mu_0 = 0.5    # internal belief
measurement_std = 0.1
sigma_s = 0.1  # variance of the sensory process this is for the next exercise

data = fep2(real_x, mu_0, measurement_std, sigma_s)
plot_variables(data)
```



Question d.4: How does the system deal with the noise? and what does σ_s mean to the agent in normal words?

[Describe your findings here](#)

It seems that noise in the sensor is dealt with pretty successfully, while internal belief noise has a larger effect on the prediction.

σ_s acts as a sort of normalising agent around the inverse model of the internal belief multiplied by the sensor input minus the forward model of the internal belief.

Question d.5: Now describe how you would implement the action in the FEP mathematical framework

[Describe it here, you can use equations](#)

An action is the minimisation of the free energy principle as a function of our internal belief and sensory reading. $a = \arg_a \min F(\mu, s(a))$

We can then use this to calculate the change in sensation with respect to our action, by applying the chain rule on $\frac{\partial F_s}{\partial a}$